

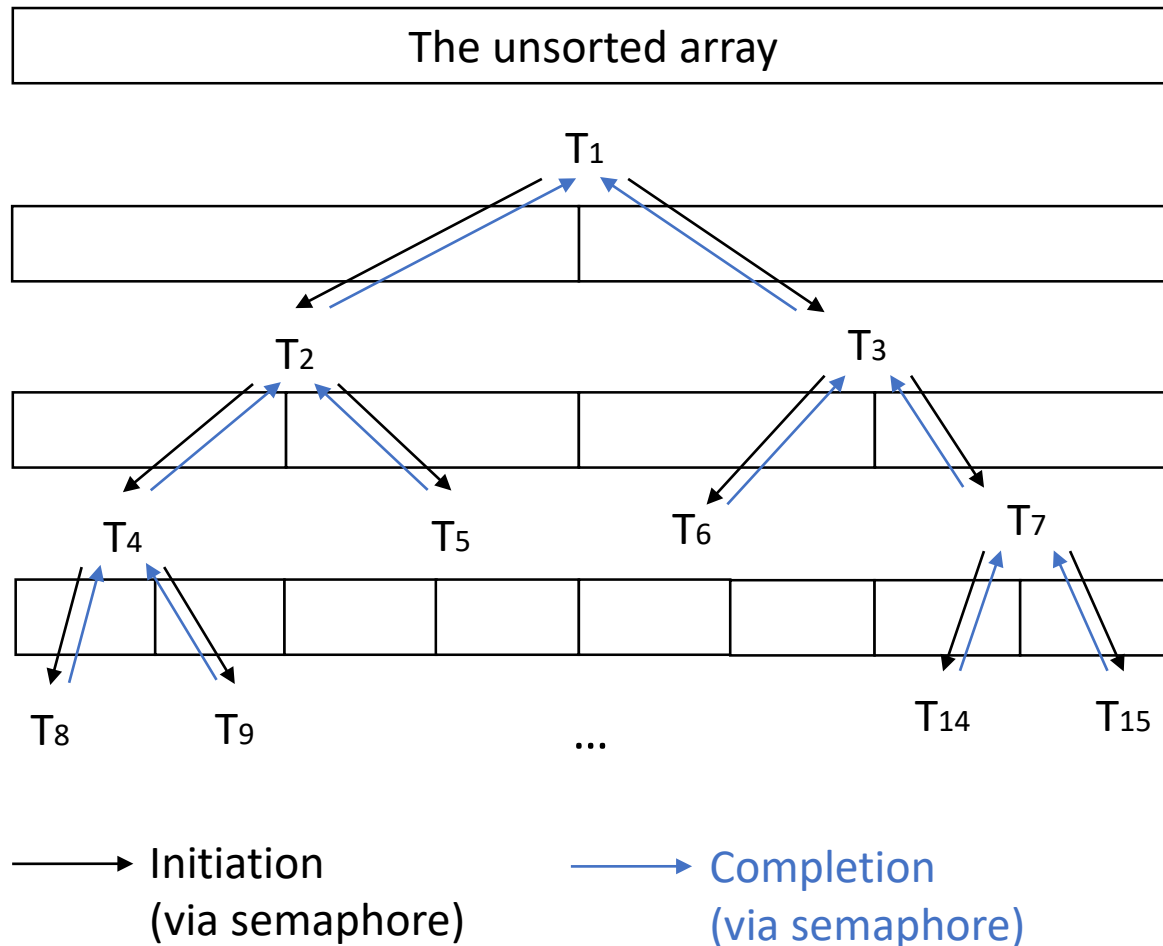
# Operating Systems Programming Assignment #3

## Parallel Merge Sort using Pthread

Prof. Li-Pin Chang

National Chiao-Tung University

# Parallel Merge Sort



- $T_1$ : the master thread
1. Divides the array into two equal sub-arrays
  2. Signals  $T_2$  and  $T_3$  (via **semaphores**) to sort the two sub-arrays
  3. Waits on  $T_2$  and  $T_3$  (via **semaphores**)
  4. Merges the two sorted sub-arrays
  5. Generate an output file

- $T_8 \sim T_{15}$ :
1. Do bubble sort on their own sub-arrays
  2. Signal their upper-level threads (via **semaphores**)

# APIs

- `<pthread.h>`

Thread management

- `Pthread_create`, `pthread_exit`

- `<semaphore.h>`

Semaphore operations

- `sem_init`, `sem_wait`, `sem_post`, `sem_getvalue`, `sem_destroy`

# Requirements

1. Prompt for the name of the input file
2. Read integers from the file
3. Do the sorting
4. Print the execution time of multi-thread sorting and single-thread sorting
  - MT sorting should be much faster than ST sorting
  - Their results must be exactly the same
5. Write the sorted array to a file
  - output1.txt → MT sorting
  - output2.txt → ST sorting

# Requirements

- The cooperation among threads must be **exactly the same** as shown in the figure
- Create all threads **in the beginning** of your program
  - Each of T1~T15 waits on its own semaphore
  - The main program signals the master thread T1 to start
  - T1 signals the 2nd-level threads T2 and T3 to start, and so on
  - Lower-level threads notify upper-level threads via semaphores; do not use `pthread_join()`
- Use **Bubble sort** at the bottom level (T8~T15)
  - For observation of speed-up

# Requirements

- Single-thread sorting
  - Use one single thread to do the same sorting, but no thread parallelism
  - 3 levels of array partitioning, bubble sort at the bottom level, and merge sub-arrays on return
  - Should be noticeably slower than the multithreaded version
- Fail to comply with the requirements will incur a score penalty
- You get 0 point if you call `qsort()` at any place in your program

# Input/output format

- Input file format:

<total # of integers><space>\n

<all integers separated by space>

- Largest input: 1,000,000 integers
- Generate your own file for testing

- Output file format:

<sorted integers separated by space>

# Header of your .c or .cpp

```
/*
```

```
Student No.: <your student id>
```

```
Student Name: <your name>
```

```
Email: <your email>
```

```
SE tag: xnxctxuxoxsx
```

```
Statement: I am fully aware that this program is not  
supposed to be posted to a public server, such as a  
public GitHub repository or a public web page.
```

```
*/
```



# Testing OS Environment

- Ubuntu 18.04
- Install as a VM or on a physical machine